

# 1 Deep Learning Theory

The following section covers the most important background theory for the experiments in Section ?? . This includes an introduction to various types of neural networks, as well as an introduction to the optimization of such networks.

## 1.1 Feedforward Neural Networks

**Feedforward neural networks** are the most basic type of neural networks. The aim of a feed-forward neural network is to approximate some function  $f^*$ , by defining a mapping  $\mathbf{y} = f(\mathbf{x}; \boldsymbol{\theta})$  and learning the parameters  $\boldsymbol{\theta}$ , that results in the best approximation of  $f^*$ . These models are called **feedforward** because there are no **feedback** connections in which the outputs of the model are fed back into itself. Instead, information flows through the function being evaluated from  $\mathbf{x}$ , through the intermediate computations used to define  $f$ , and finally to the output  $\mathbf{y}$ . Feedforward neural networks generally consists of multiple **layers**, arranged in a chain structure, with each layer being a function of the layer that preceded it [2].

## 1.2 Fully-connected Layers

The most simple type of layer found in a feedforward neural network is the **fully-connected layer**. The fully-connected layer usually consists of some learnable parameter matrix  $\mathbf{W}$  and learnable parameter vector  $\mathbf{b}$ , as well as a non-linear **activation function**  $g$  (which will be covered further in Section 1.6.1). In this case, the  $i$ 'th layer is defined as [2]

$$\mathbf{h}^{(i)} = \begin{cases} g^{(i)} (\mathbf{W}^{(i)\top} \mathbf{h}^{(i)} + \mathbf{b}^{(i)}) & \text{if } i > 1 \\ g^{(1)} (\mathbf{W}^{(1)\top} \mathbf{x} + \mathbf{b}^{(1)}) & \text{if } i = 1 \end{cases} \quad (1)$$

## 1.3 Convolutional Layer

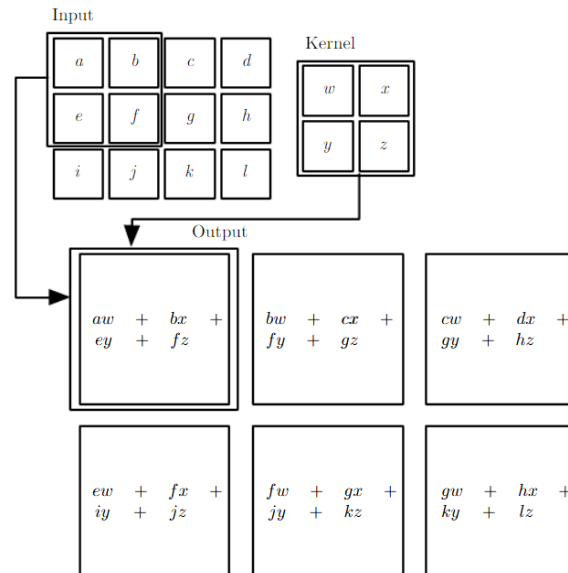


Figure 1: An example of applying a 2d kernel on an input [2].

A **convolutional layer** is a specialized kind of feedforward layer, usually used in analysis of time-series or image data [2]. If a network has at least one convolutional layer, it is called a **Convolutional neural network (CNN)**.

The convolutional layer consists of a set of **kernels**, each to be applied to the entire input vector, where each kernel is a learnable parameter matrix  $k \times k$  [3]. Each kernel is applied on the input to produce a **feature map**. The kernels are applied to the input by "sliding" over the input (where the step size is called **stride**). Each  $k \times k$  grid of the input is then used to compute the dot-product between the grid and each kernel, which is then placed in the corresponding feature map of each kernel, as visualized in Figure 1. [1]. To control the dimensions of the output, one might **pad** the sides with a constant value. Commonly, zero is used as the padding-value.

As seen in Figure 1, each kernel produces a linear combination of all pixel values in a neighbourhood defined by the size of the kernel. Thus, unlike a fully-connected layer, a convolutional layer captures the high correlation between a pixel and its neighbours. Further, by limiting the size of the kernel, the network will use much fewer parameters, than if a fully-connected layer would be used instead [2].

## 1.4 Recurrent Neural Networks

### 1.4.1 Long Short-Term Memory Unit

- Convolutional LSTM

### 1.4.2 Gated Recurrent Unit

- Convolutional GRU

## 1.5 Transformer

## 1.6 Training a Neural Network

### 1.6.1 Activation function